




11 DECEMBER 2022

Introduction to system/backend programming

System programming

System programming?

- System programming is about **mechanical sympathy**
- It's about knowing **how and why** stuff work the way they do
- And most importantly – being able to critically think about your program or **system**



Copy
pasting code
from
stackoverflow



Being
a system
programmer



Array copy

```
void copyij(long int src[2048][2048], long int dst[2048][2048])
{
    long int i,j;
    for (i = 0; i < 2048; i++)
        for (j = 0; j < 2048; j++)
            dst[i][j] = src[i][j];
}
```

```
void copyji(long int src[2048][2048], long int dst[2048][2048])
{
    long int i,j;
    for (j = 0; j < 2048; j++)
        for (i = 0; i < 2048; i++)
            dst[i][j] = src[i][j];
}
```



Array copy

```
void copyij(long int src[2048][2048], long int dst[2048][2048])
{
    long int i,j;
    for (i = 0; i < 2048; i++)
        for (j = 0; j < 2048; j++)
            dst[i][j] = src[i][j];
}

void copyji(long int src[2048][2048], long int dst[2048][2048])
{
    long int i,j;
    for (j = 0; j < 2048; j++)
        for (i = 0; i < 2048; i++)
            dst[i][j] = src[i][j];
}
```

copyij time spent: 0.016467 s (16ms)
copyji time spent: 0.039325 s (39ms) **2x slowdown**



Array copy

```
void copyij(long int src[2048][2048], long int dst[2048][2048])
{
    long int i,j;
    for (i = 0; i < 2048; i++)
        for (j = 0; j < 2048; j++)
            dst[i][j] = src[i][j];
}

void copyji(long int src[2048][2048], long int dst[2048][2048])
{
    long int i,j;
    for (j = 0; j < 2048; j++)
        for (i = 0; i < 2048; i++)
            dst[i][j] = src[i][j];
}
```

copyij time spent: 0.016467 s (16ms)
copyji time spent: 0.039325 s (39ms) x2
slowdown

Row-major

0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

Column-major

0	4	8	12
1	5	9	13
2	6	10	14
3	7	11	15

0-3	Node 1
4-7	Node 2
8-11	Node 3
12-15	Node 4



System calls

- Computer programs (usually) run within the confines of the OS/kernel



System calls

- Computer programs (usually) run within the confines of the OS
- In order for them to interface with the external world (R/W files, R/W on the network) they need the help of the OS



System calls

- Computer programs (usually) run within the confines of the OS
- In order for them to interface with the external world (R/W files, R/W on the network) they need the help of the OS
- This is achieved via System Calls (or syscalls)
 - That's the interface the operating system (the kernel) gives to programs
 - This allows to give controlled access to the hardware (hdd, network card).

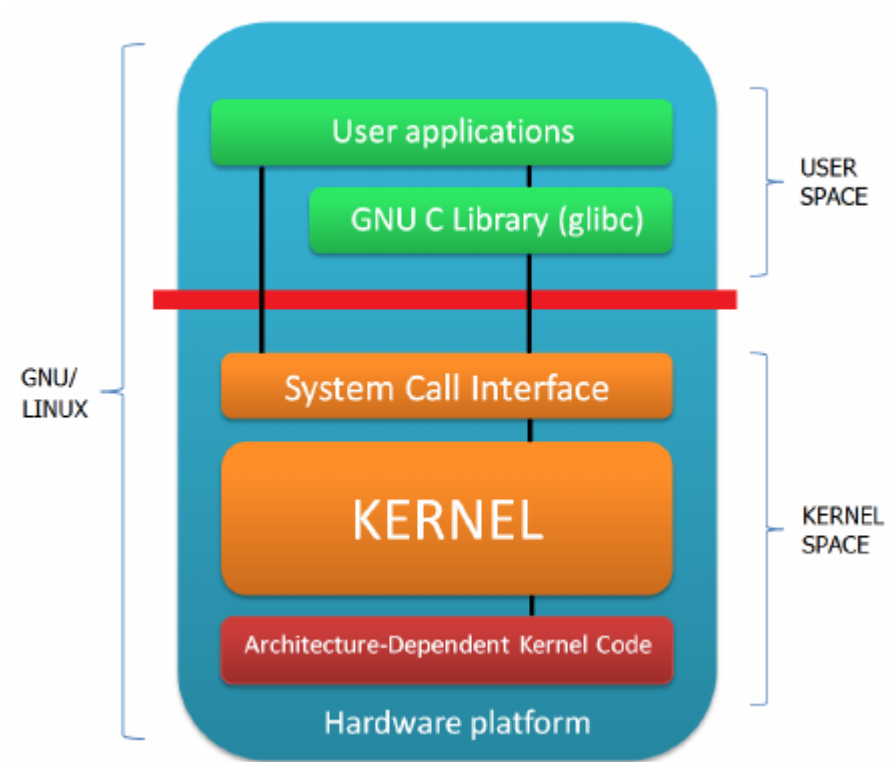


System calls

- Computer programs (usually) run within the confines of the OS
- In order for them to interface with the external world (R/W files, R/W on the network) they need the help of the OS
- This is achieved via System Calls (or syscalls)
 - That's the interface the operating system (the kernel) gives to programs
 - This allows to give controlled access to the hardware (hdd, network card).
- System calls are generally implemented as an **exception handler** i.e they cause the current **execution context** to jump to the **kernel** and then get routed to the correct function
 - When this happens the kernel performs some action on behalf of the user



System calls



System calls (2)

```
with open("test.txt", "w+") as f:  
    f.write("Hello world")  
    f.seek(0)  
    print(f.read())
```

```
# strace python3 filewrite.py  
openat(AT_FDCWD, "test.txt", O_RDWR|O_CREAT|O_TRUNC|  
O_CLOEXEC, 0666) = 3  
<omitted for brevity>  
write(3, "Hello world", 11)          = 11  
lseek(3, 0, SEEK_SET)                = 0  
read(3, "Hello world", 12)          = 11  
write(1, "Hello world\n", 12Hello world) = 12
```

<https://man7.org/linux/man-pages/man2/open.2.html>

<https://man7.org/linux/man-pages/man2/write.2.html>

<https://man7.org/linux/man-pages/man2/lseek.2.html>

<https://man7.org/linux/man-pages/man2/read.2.html>



Operating System

- We saw the term OS and you've probably heard of Linux/Windows/Mac but...



Operating System

- We saw the term OS and you've probably heard of Linux/Windows/Mac but...
 - Have you asked yourself **what exactly IS the OS?**



Operating System

- We saw the term OS and you've probably heard of Linux/Windows/Mac but...
 - Have you asked yourself **what exactly IS the OS?**
 - Is it the taskbar or the task manager on Windows?
 - Is it the root shell on a Linux system?



Operating System

- We saw the term OS and you've probably heard of Linux/Windows/Mac but...
 - Have you asked yourself **what exactly IS the OS?**
 - Is it the taskbar or the task manager on Windows?
 - Is it the root shell on a Linux system?
- The answer is all **NO**



Operating System

- We saw the term OS and you've probably heard of Linux/Windows/Mac but...
 - Have you asked yourself **what exactly IS the OS?**
 - Is it the taskbar or the task manager on Windows?
 - Is it the root shell on a Linux system?
- The answer is all **NO**
- The OS is really a bunch of **exception handlers** (like the syscall handler) + accompanying code like:
 - Architecture support code
 - The process scheduler
 - The virtual memory subsystem
 - Hardware device drivers
 - Filesystems



Operating System

- The OS is always running as part of every ordinary process
 - Thus we conclude each process has a user space portion and kernel (privileged) portion
 - Kernel is not copied, the state is shared I.e currently running process or misc. accounting
 - There's a constant jump between user space/kernel space via (traps) syscalls and asynchronous exceptions (interrupts)



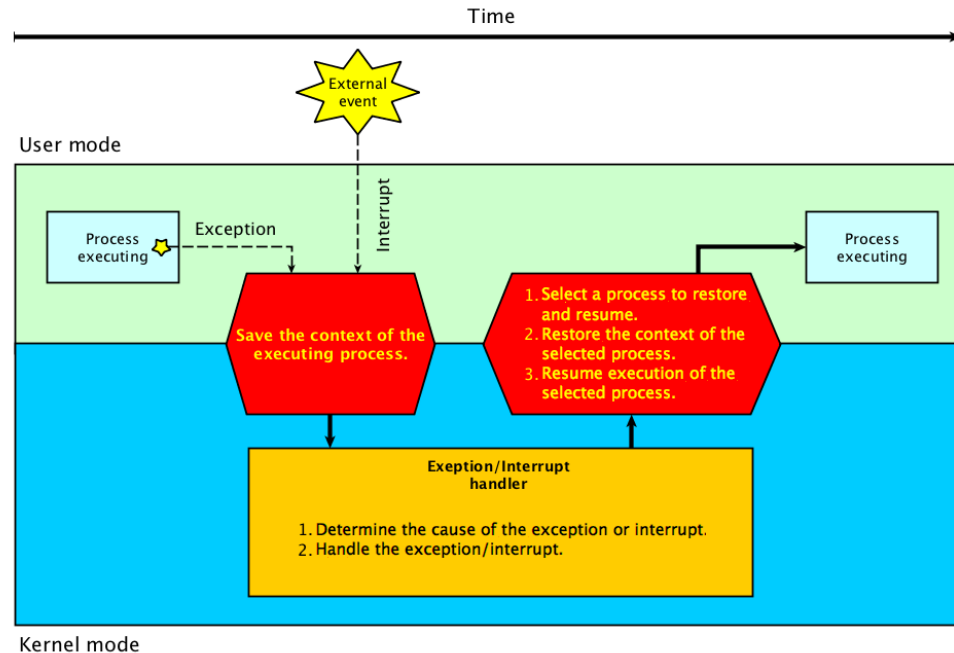
Operating System (2)

- The OS is always running as part of every ordinary process
 - Thus we conclude each process has a user space portion and kernel (privileged) portion
 - Kernel is not copied, the state is shared I.e currently running process or misc. accounting
 - There's a constant jump between user space/kernel space via (traps) syscalls and asynchronous exceptions (interrupts)
- Asynchronous Exception (interrupts)
 - Causes current execution to suspend and jump to respective handler, handle the event and return back to the user process.



Operating System (2)

- The OS is always running as part of every ordinary process
 - Thus we conclude each process has a user space portion and kernel (privileged) portion
 - Kernel is not copied, the state is shared I.e currently running process or misc. accounting
 - There's a constant jump between user space/kernel space via (traps) syscalls and asynchronous exceptions (interrupts)



Operating System (2)

- The OS is always running as part of every ordinary process
 - Thus we conclude each process has a user space portion and kernel (privileged) portion
 - Kernel is not copied, the state is shared I.e currently running process or misc. accounting
 - There's a constant jump between user space/kernel space via (traps) syscalls and asynchronous exceptions (interrupts)
- Asynchronous Exception (interrupts)
 - Causes current execution to suspend and jump to respective handler, handle the event and return back to the user process.
 - Type of interrupts
 - Keyboard key stroke
 - Network packet arrival
 - hdd/ssd byte transferred
 - Time interrupt




Operating System (2)

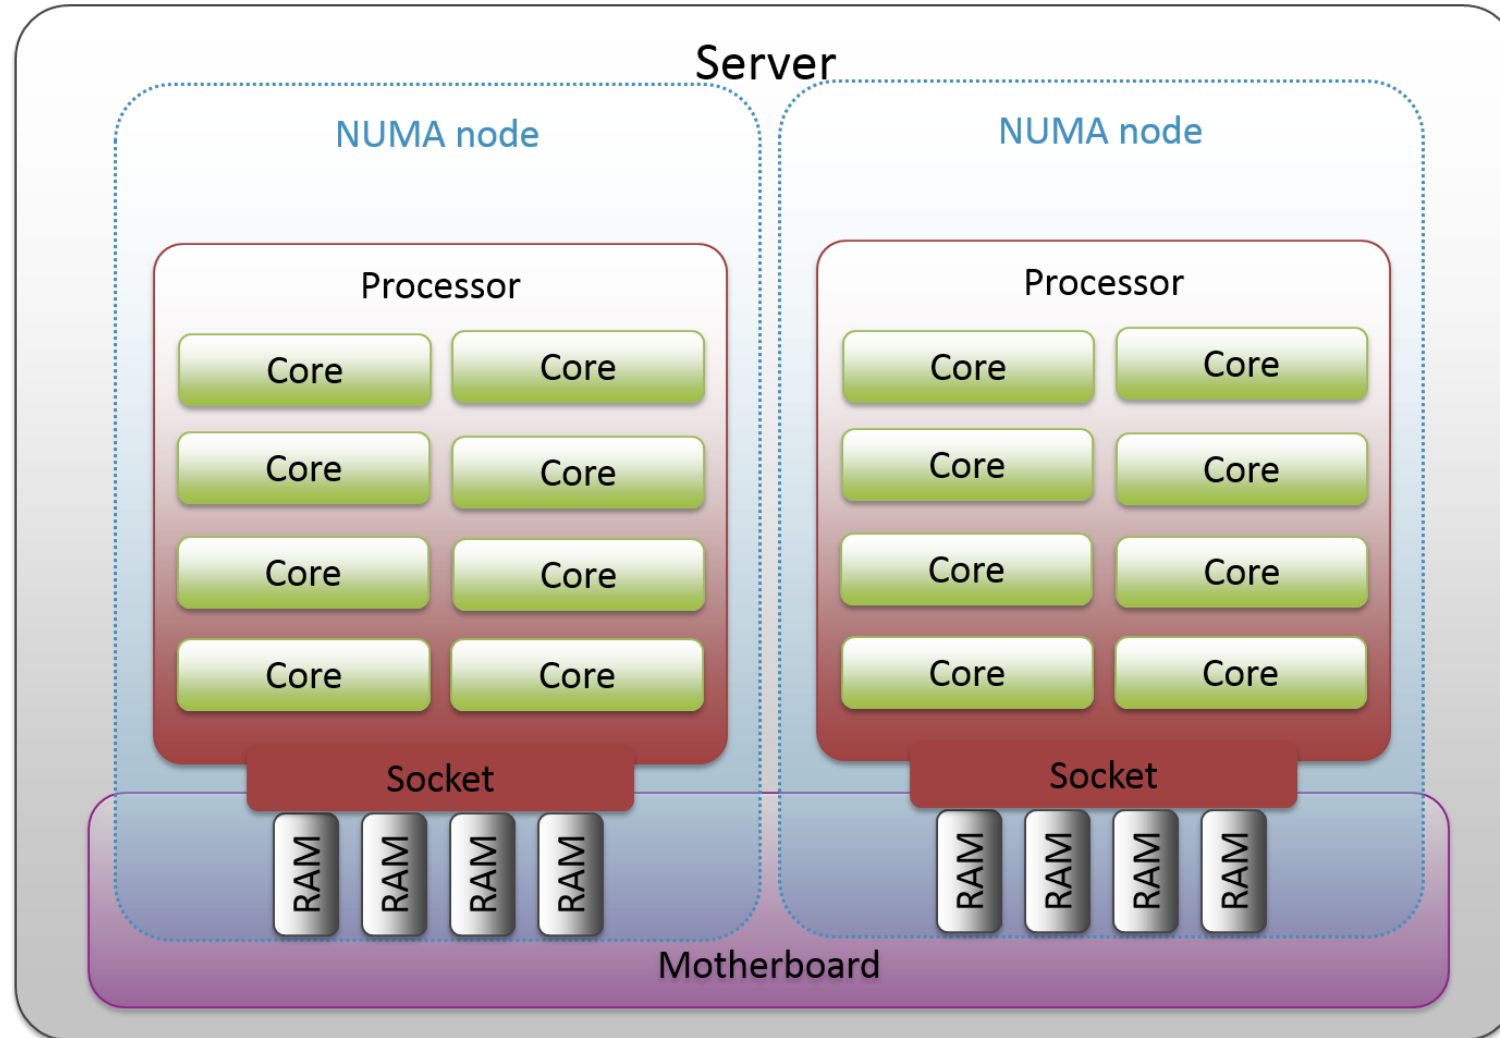
- The OS is always running as part of every ordinary process
 - Thus we conclude each process has a user space portion and kernel (privileged) portion
 - Kernel is not copied, the state is shared i.e currently running process or misc. accounting
 - There's a constant jump between user space/kernel space via (traps) syscalls and asynchronous exceptions (interrupts)
- Asynchronous Exception (interrupts)
 - Causes current execution to suspend and jump to respective handler, handle the event and return back to the user process.
 - Type of interrupts
 - Keyboard key stroke
 - Network packet arrival
 - hdd/ssd byte transferred
 - Time interrupt
- There are 2 other types of exceptions – faults and abort
 - Faults – Unintentional but recoverable e.g. page faults (recoverable), protection fault (unrecoverable).
 - Aborts – similar to unrecoverable faults, causes program to terminate i.e divide-by-zero, parity error, generally signals some hardware problem in the CPU.



Non-uniform memory access



Non-uniform memory access



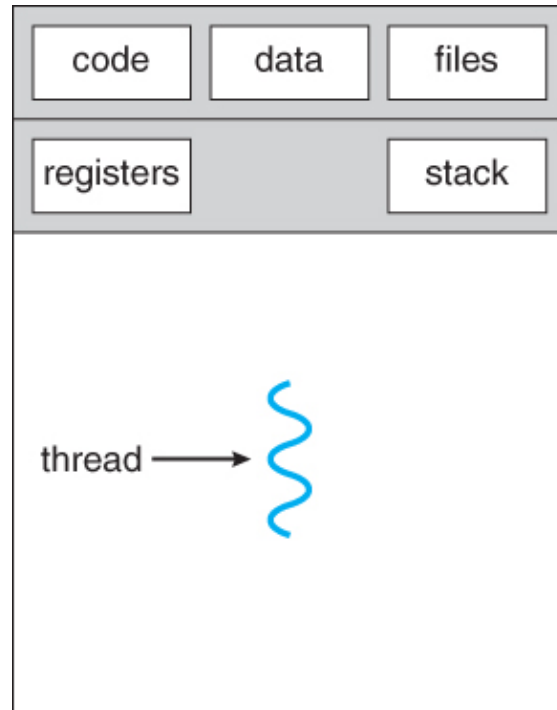
Process execution

- Process is an instance of a running program
 - It contains all the state necessary for the process to run – opened files, code, data, CPU state (program counter, stack, registers)
 - At least one thread of execution




Process execution

- Process is an instance of a running program
 - It contains all the state necessary for the process to run – opened files, code, data, CPU state (program counter, stack, registers), memory address space
 - At least one thread of execution




single-threaded process



multithreaded process



Concurrency vs Parallelism





Thank you

For more information, contact SUSE at:

+1 800 796 3700 (U.S./Canada)

Maxfeldstrasse 5

90409 Nuremberg

www.suse.com

© 2022 SUSE LLC. All Rights Reserved.
SUSE and the SUSE logo are registered
trademarks of SUSE LLC in the United States
and other countries. All third-party
trademarks are the property of their
respective owners.